

Project

Manasa

2024-03-17

R Markdown

```
if (!requireNamespace("readr", quietly = TRUE)) install.packages("readr")
if (!requireNamespace("dplyr", quietly = TRUE)) install.packages("dplyr")
if (!requireNamespace("tidyverse", quietly = TRUE)) install.packages("tidyverse")
if (!requireNamespace("lubridate", quietly = TRUE)) install.packages("lubridate")

#Load the libraries
library(lubridate)

## 
## Attaching package: 'lubridate'

## The following objects are masked from 'package:base':
## 
##     date, intersect, setdiff, union

library(readr)
library(ggplot2)
library(dplyr)

## 
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
## 
##     filter, lag

## The following objects are masked from 'package:base':
## 
##     intersect, setdiff, setequal, union

library(tidyverse)

crd <- read_csv("/Users/manasavishnumoorthy/Desktop/IDMP/CRD.csv",
                na = c("", "NA", " "), show_col_types = FALSE)

incident <- read_csv("/Users/manasavishnumoorthy/Desktop/IDMP/incidents_data.csv",
```

```

        na = c("", "NA", " "), show_col_types = FALSE)

non_motorists <- read_csv("/Users/manasavishnumoorthy/Desktop>IDMP/non_motorist_data.csv",
                           na = c("", "NA", " "), show_col_types = FALSE)

# Convert 'Local Case Number' to character in both data frames
crd$`Local Case Number` <- as.character(crd$`Local Case Number`)
non_motorists$`Local Case Number` <- as.character(non_motorists$`Local Case Number`)
incident$`Local Case Number` <- as.character(incident$`Local Case Number`)

combined_data <- crd %>%
  full_join(non_motorists, by = c("Report Number", "Local Case Number")) %>%
  full_join(incident, by = c("Report Number", "Local Case Number"))

# Check the first few rows of the combined data
#str(combined_data)

```

Data preprocessing Standardize text variables: Convert all character variables to lowercase for consistency.

```

combined_data <- combined_data %>%
  mutate(across(where(is.character), tolower))

```

Handle missing values: Replace missing values in character columns with “unknown” and in numeric columns with 0 (or another appropriate value).

```

combined_data <- combined_data %>%
  mutate(across(where(is.character), ~replace_na(., "unknown"))) %>%
  mutate(across(where(is.numeric), ~replace_na(., 0)))

```

Create new variables from ‘Crash Date/Time’: Parse the ‘Crash Date/Time’ and extract useful components such as the date, hour, day of the week, etc.

```

# This assumes that the date/time column is named 'Crash Date/Time.x' in your combined_data
combined_data <- combined_data %>%
  mutate(`Crash Date/Time` = mdy_hms(`Crash Date/Time.x`)) # Convert to date-time object

# Extract components
combined_data <- combined_data %>%
  mutate(
    Date = as.Date(`Crash Date/Time`), # Extract date
    Hour = hour(`Crash Date/Time`), # Extract hour
    Weekday = wday(`Crash Date/Time`, label = TRUE), # Extract day of the week
    Month = month(`Crash Date/Time`, label = TRUE) # Extract month
  )

#str(combined_data)

```

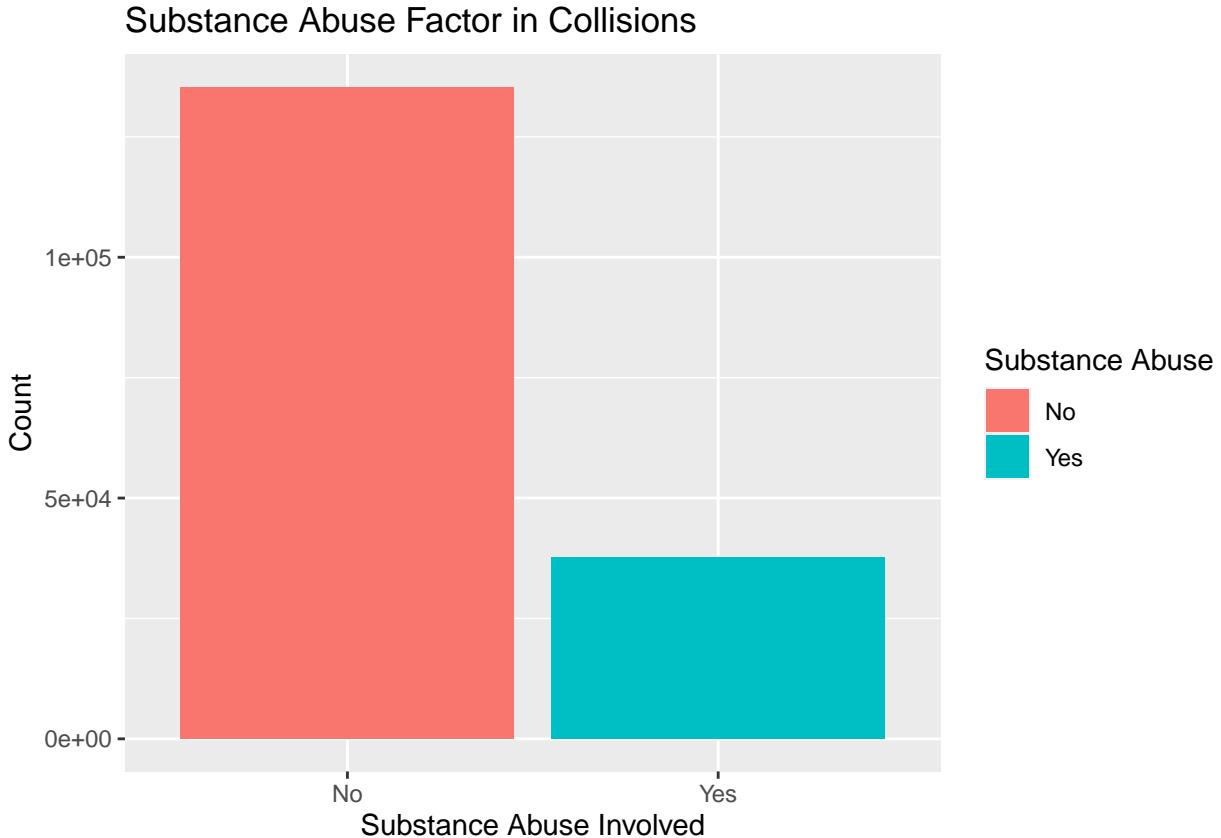
Substance Abuse Factor: Investigating the factor of substance abuse in collisions.

```

combined_data$`Substance Abuse` <- ifelse(combined_data$`Driver Substance Abuse.x` != "none detected" &
                                         combined_data$`Driver Substance Abuse.x` != "unknown", "Yes",
                                         "No")

ggplot(combined_data, aes(x = `Substance Abuse`, fill = `Substance Abuse`)) +
  geom_bar() +
  labs(title = "Substance Abuse Factor in Collisions", x = "Substance Abuse Involved", y = "Count")

```



Accidents by Light Conditions and Time of Day: It might be interesting to see how light conditions impact accidents, and if there are peak times.

```

library(ggplot2)
library(dplyr)

# Filter out unwanted 'Light.x' categories
filtered_data <- combined_data %>%
  filter(!Light.x %in% c("n/a", "other", "unknown"))

# Count the number of accidents by light condition
light_data <- filtered_data %>%
  count(Light.x) %>%
  rename(Number_of_Accidents = n)

# Create a pie chart
ggplot(light_data, aes(x = "", y = Number_of_Accidents, fill = Light.x)) +
  geom_bar(stat = "identity", width = 1) +
  coord_polar(theta = "y") # This makes it a pie chart

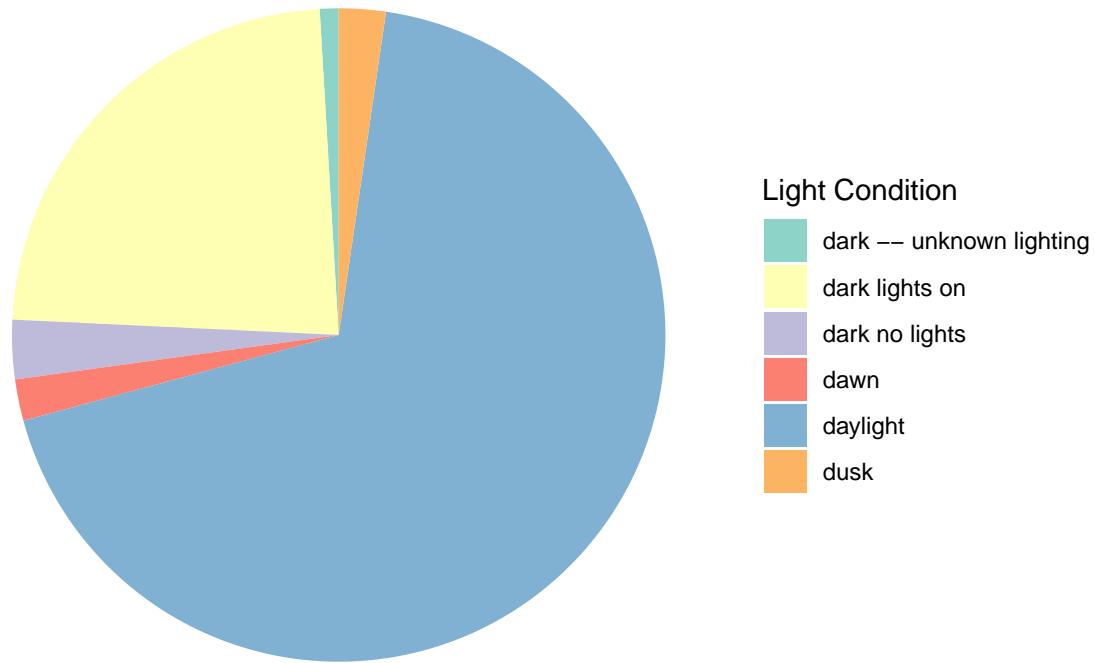
```

```

scale_fill_brewer(palette = "Set3") +
labs(title = "Proportion of Accidents by Light Conditions", x = "", y = "", fill = "Light Condition")
theme_void() # This removes the background, gridlines, and text

```

Proportion of Accidents by Light Conditions



Heatmap of Collisions by Hour and Day of the Week This visualization will help us understand when most collisions occur during the week.

```

library(ggplot2)
library(dplyr)

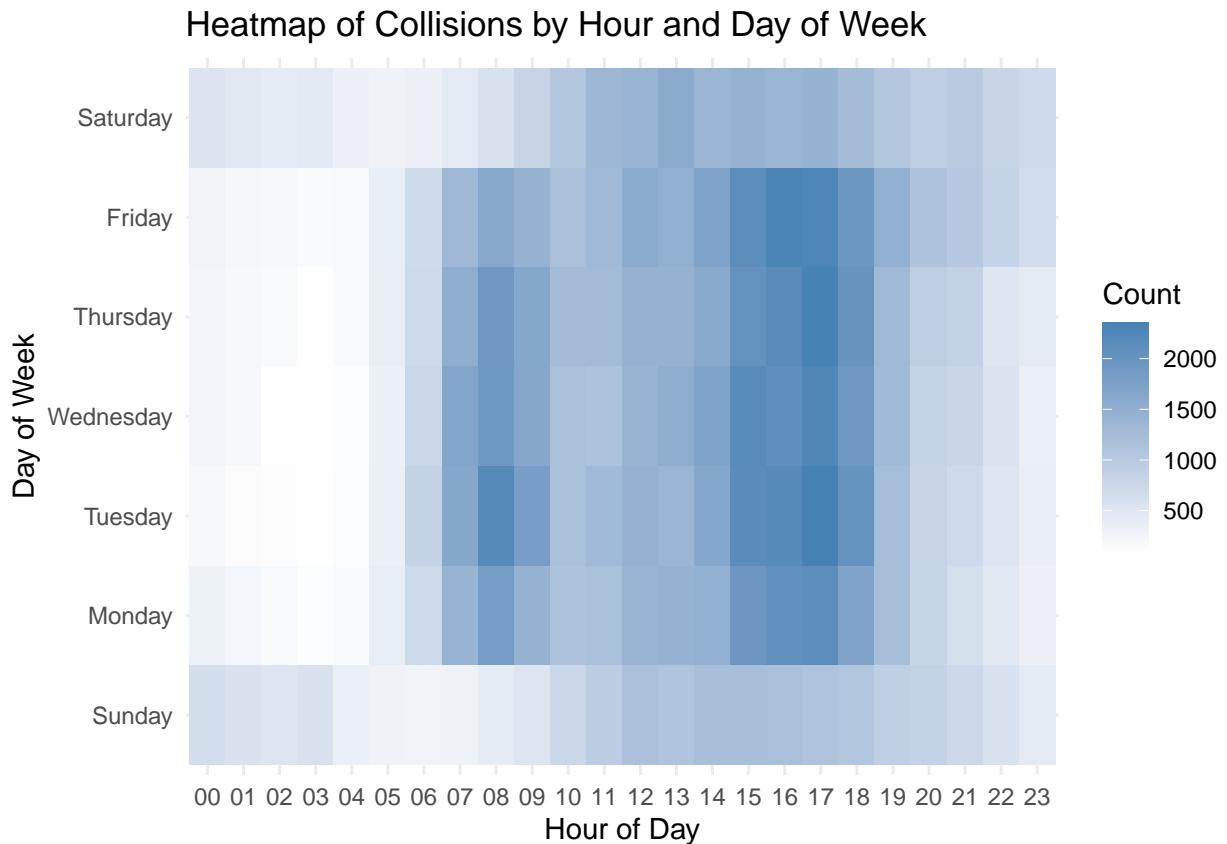
# Create a column for the day of the week and hour
combined_data$DayOfWeek <- weekdays(as.Date(combined_data$`Crash Date/Time`))
combined_data$Hour <- format(combined_data$`Crash Date/Time` , "%H")

# Prepare the data for the heatmap
heatmap_data <- combined_data %>%
  filter(!is.na(DayOfWeek) & DayOfWeek != "na") %>% # Filter out the "NA" or "na" entries
  group_by(DayOfWeek, Hour) %>%
  summarize(Count = n(), .groups = 'drop') %>%
  mutate(DayOfWeek = factor(DayOfWeek, levels = c("Sunday", "Monday", "Tuesday", "Wednesday", "Thursday",
                                                 "Friday", "Saturday")))

# Create the heatmap
ggplot(heatmap_data, aes(x = Hour, y = DayOfWeek, fill = Count)) +
  geom_tile() +
  scale_fill_gradient(low = "white", high = "steelblue") +
  labs(title = "Heatmap of Collisions by Hour and Day of Week", x = "Hour of Day", y = "Day of Week")

```

```
theme_minimal()
```

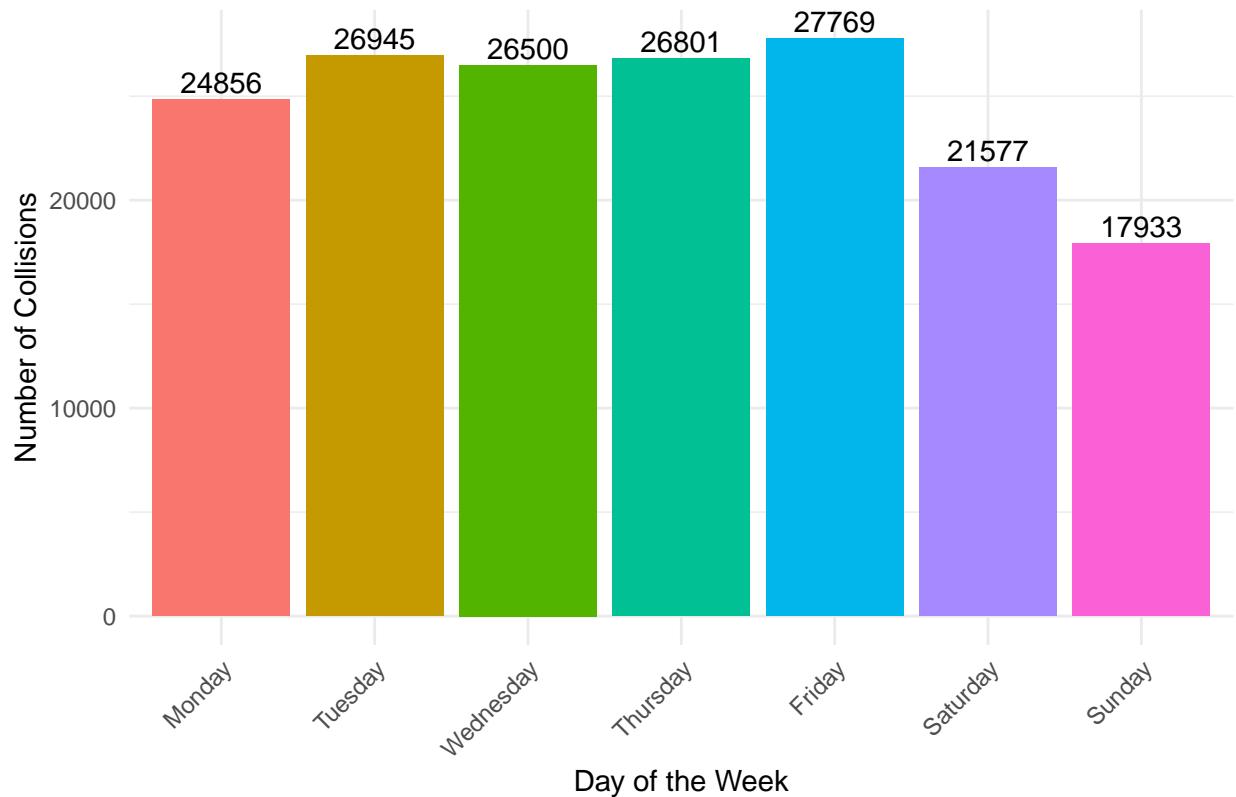


```
library(ggplot2)
library(dplyr)

# Assuming the 'DayOfWeek' variable is already created and 'NA' is one of the levels
collisions_by_day_filtered <- combined_data %>%
  filter(is.na(DayOfWeek) & DayOfWeek != "na") %>% # Filter out the "NA" or "na" entries
  group_by(DayOfWeek) %>%
  summarise(NumberOfCollisions = n(), .groups = 'drop') %>%
  mutate(DayOfWeek = factor(DayOfWeek, levels = c("Monday", "Tuesday", "Wednesday", "Thursday", "Friday",
                                                 "Saturday", "Sunday")))

# Now, create the bar chart without the "NA" category
ggplot(collisions_by_day_filtered, aes(x = DayOfWeek, y = NumberOfCollisions, fill = DayOfWeek)) +
  geom_bar(stat = "identity", show.legend = FALSE) +
  geom_text(aes(label = NumberOfCollisions), vjust = -0.3) +
  labs(title = "Number of Collisions by Day of the Week", x = "Day of the Week", y = "Number of Collisions")
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

Number of Collisions by Day of the Week



ARIMA

```

library(dplyr)
library(lubridate)
library(forecast)

## Registered S3 method overwritten by 'quantmod':
##   method           from
##   as.zoo.data.frame zoo

library(tibble)
library(ggplot2)

# Convert the Crash Date/Time to a proper Date format (assuming it is in a recognisable format)
combined_data$Crash_Date <- as.Date(mdy_hms(combined_data$`Crash Date/Time.x`))

# Remove rows with NA in Crash_Date (if any conversion failed)
combined_data <- combined_data %>%
  filter(!is.na(Crash_Date))

# Aggregate data by Crash_Date to get the number of collisions per day
daily_collisions <- combined_data %>%
  group_by(Crash_Date) %>%
  summarise(Number_of_Collisions = n(), .groups = 'drop')

```

```

# Generate the time series object from daily collision counts
collision_ts <- ts(daily_collisions$Number_of_Collisions, frequency = 365)

# Fit the ARIMA model
fit <- auto.arima(collision_ts)

# Check model summary
summary(fit)

## Series: collision_ts
## ARIMA(5,1,2)
##
## Coefficients:
##             ar1      ar2      ar3      ar4      ar5      ma1      ma2
##            0.0974 -0.3863 -0.2462 -0.2249 -0.2886 -0.8447  0.2640
## s.e.    0.0610  0.0696  0.0349  0.0291  0.0324  0.0535  0.1214
##
## sigma^2 = 226.5:  log likelihood = -13573.89
## AIC=27163.78   AICc=27163.82   BIC=27212.56
##
## Training set error measures:
##           ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -0.02631048 15.03234 11.64846 -9.676269 26.68143 0.6615329
##           ACF1
## Training set -0.01223579

library(ggplot2)
library(dplyr)
library(lubridate)

# Assuming 'daily_collisions' is your data frame with the 'Crash_Date' and 'Number_of_Collisions' column

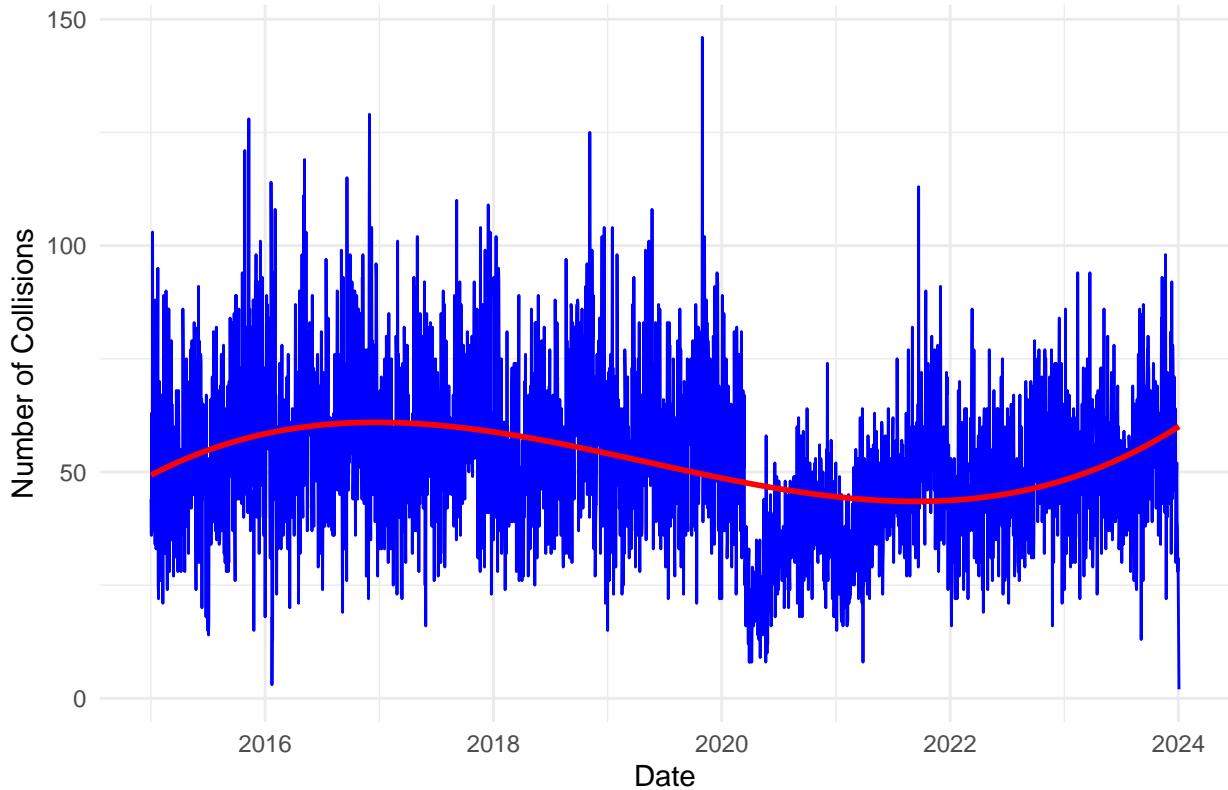
# Ensure that the dates are in the proper Date format
daily_collisions$Crash_Date <- as.Date(daily_collisions$Crash_Date)

#Plot using ggplot2
#ggplot(daily_collisions, aes(x = Crash_Date, y = Number_of_Collisions)) +
#  geom_line(color = "blue") +
#  geom_smooth(method = "lm", color = "red", se = FALSE) + # Trend line
#  labs(title = "Historical Collision Data",
#       x = "Date",
#       y = "Number of Collisions") +
#  theme_minimal()
# 

ggplot(daily_collisions, aes(x = Crash_Date, y = Number_of_Collisions)) +
  geom_line(color = "blue") +
  geom_smooth(method = "lm", formula = y ~ poly(x, 3), color = "red", se = FALSE) + # Quadratic trend line
  labs(title = "Historical Collision Data with Non-linear Trend",
       x = "Date",
       y = "Number of Collisions") +
  theme_minimal()

```

Historical Collision Data with Non-linear Trend



```
library(forecast)
library(ggplot2)

# Assuming 'fit' is your previously fitted ARIMA model
# and 'daily_collisions' is the data frame with historical collision data

# Forecast the next 30 days using the fitted ARIMA model
forecast_future <- forecast(fit, h = 30)

# Create a data frame for the forecast dates and predictions
last_date <- max(daily_collisions$Crash_Date) # Last date in your historical data
forecast_dates <- seq.Date(from = last_date + 1, by = "day", length.out = 30)

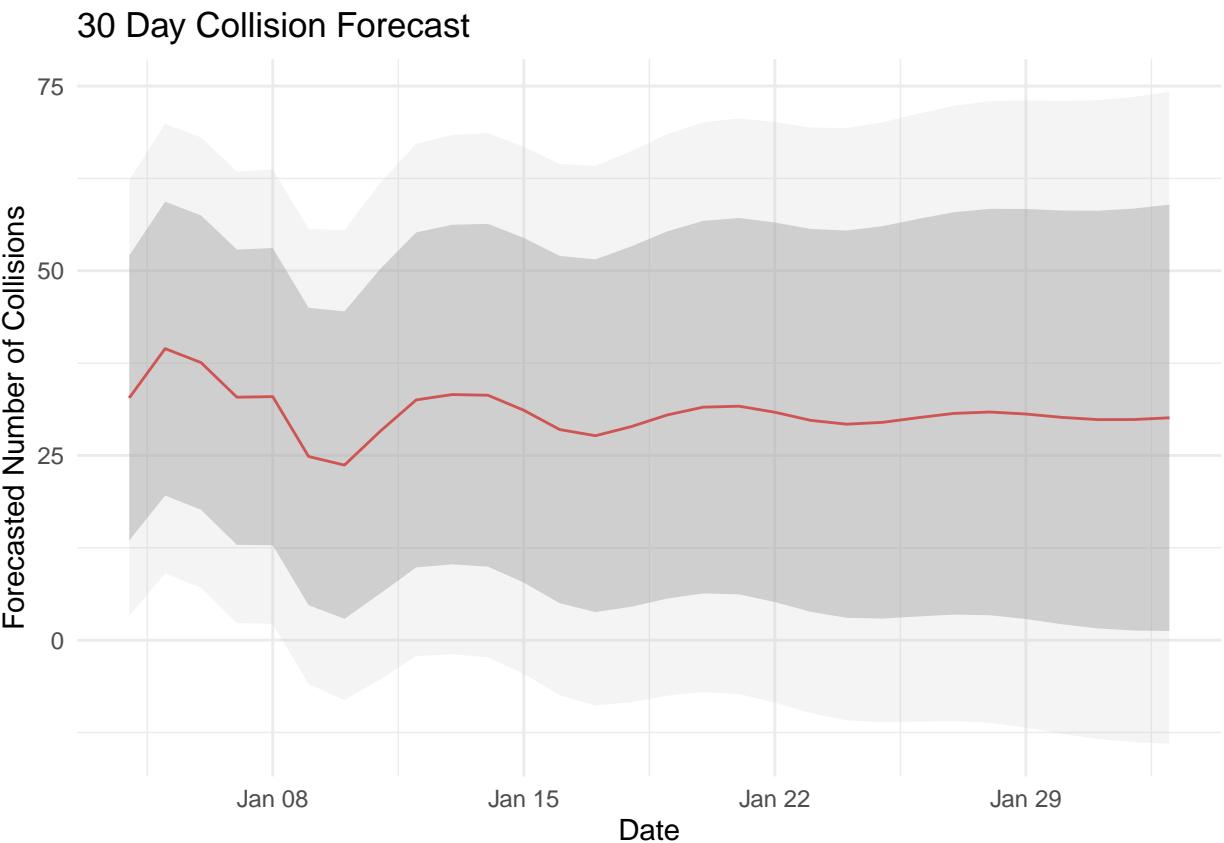
forecast_data <- data.frame(
  Date = forecast_dates,
  Forecast = as.numeric(forecast_future$mean),
  Lower_80 = forecast_future$lower[, "80%"],
  Upper_80 = forecast_future$upper[, "80%"],
  Lower_95 = forecast_future$lower[, "95%"],
  Upper_95 = forecast_future$upper[, "95%"]
)

# Plot the forecast using ggplot2
ggplot(forecast_data, aes(x = Date)) +
  geom_line(aes(y = Forecast), color = "red") +
```

```

geom_ribbon(aes(ymin = Lower_95, ymax = Upper_95), fill = "grey80", alpha = 0.2) +
  geom_ribbon(aes(ymin = Lower_80, ymax = Upper_80), fill = "grey60", alpha = 0.4) +
  labs(title = "30 Day Collision Forecast",
       x = "Date",
       y = "Forecasted Number of Collisions") +
  theme_minimal()

```

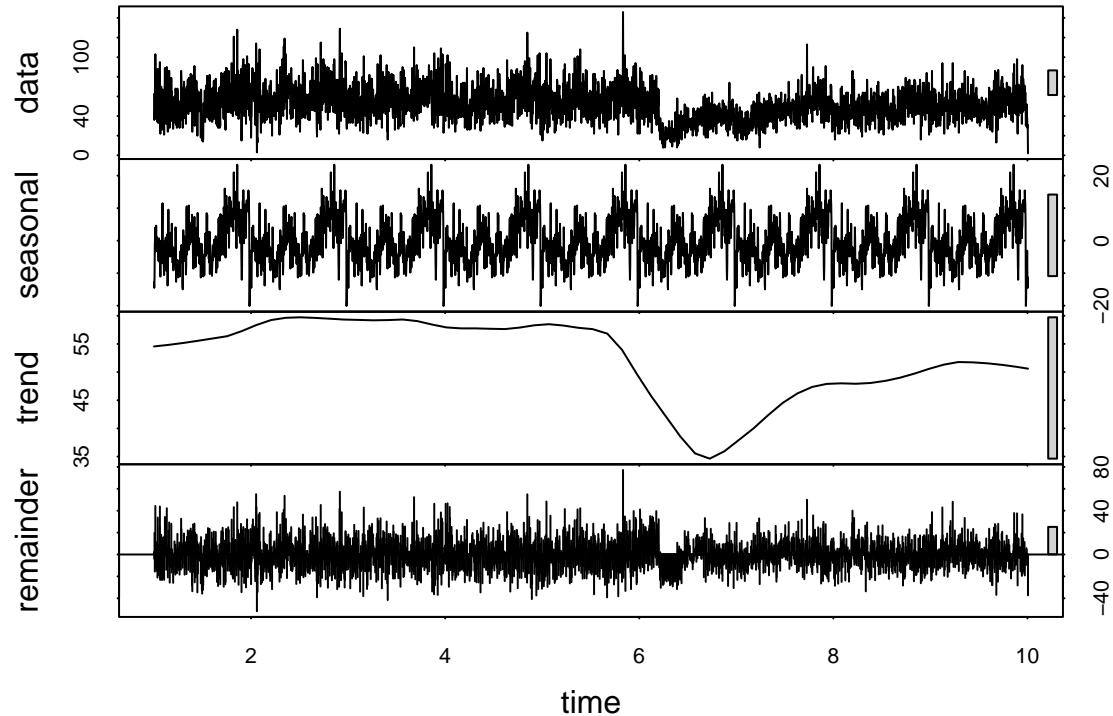


```

# Time series decomposition
decomposition <- stl(collision_ts, s.window = "periodic")

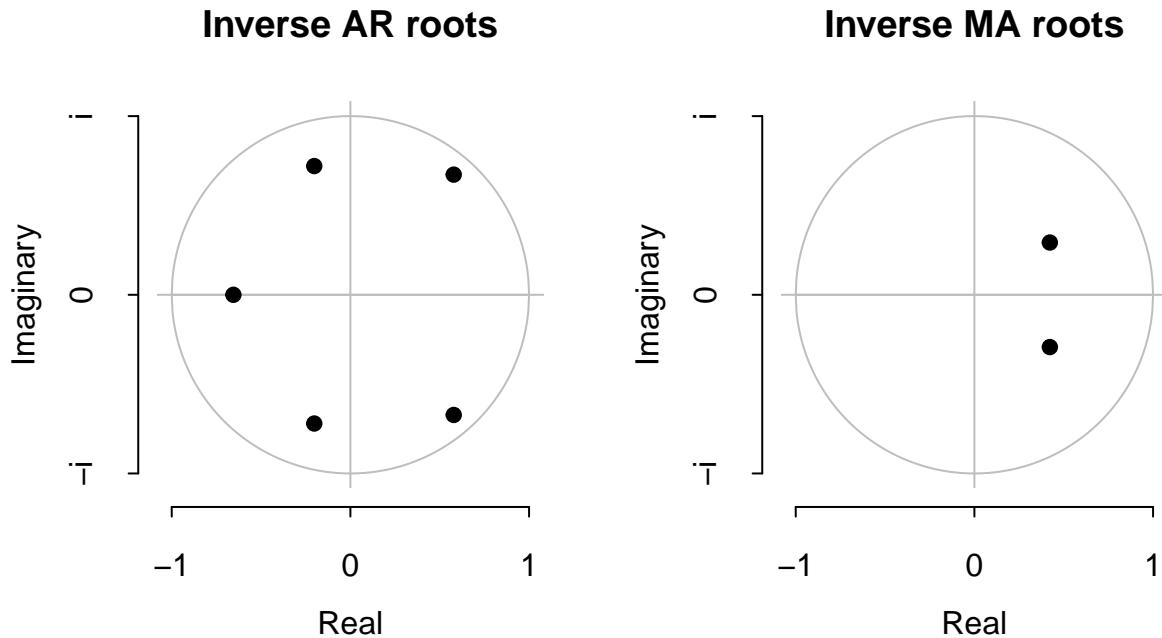
# Plot decomposed components
plot(decomposition)

```



```
# Model selection
# Example: Fit a SARIMA model
sarima_model <- auto.arima(collision_ts, seasonal = TRUE)

# Model validation
# Diagnostic plots
plot(sarima_model)
```



```
# Forecasting
forecast_values <- forecast(sarima_model, h = 30)

# Evaluate forecast accuracy
accuracy(forecast_values)
```

```
##               ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -0.02631048 15.03234 11.64846 -9.676269 26.68143 0.6615329
##                               ACF1
## Training set -0.01223579
```

Regression Analysis

```
# Load required libraries
library(ggplot2)
library(dplyr)
library(tidyr)
library(caret)

## Loading required package: lattice

# Assuming 'Injury_Severity.x' is a factor variable indicating the severity of crashes
# Let's encode it into numeric values for regression analysis
combined_data$Injury_Severity_Num <- as.numeric(factor(combined_data$`Injury Severity.x`))
```

```

# Selecting different variables for the regression model
model_data <- combined_data %>%
  select(`Injury_Severity_Num`, `Weather`, `Surface Condition`, `Light`, `Vehicle Year`, `Vehicle Make`)

# Convert Weather, Surface_Condition, Light, Vehicle_Make, and Vehicle_Model to factors
model_data$Weather <- as.factor(model_data$Weather)
model_data$Surface_Condition <- as.factor(model_data$`Surface Condition`)
model_data$Light <- as.factor(model_data$Light)
model_data$Vehicle_Year <- as.factor(model_data$`Vehicle Year`)
model_data$Vehicle_Make <- as.factor(model_data$`Vehicle Make`)
model_data$Vehicle_Model <- as.factor(model_data$`Vehicle Model`)

# Encoding factors into numeric values for regression analysis
model_data <- model_data %>%
  mutate(Weather_Num = as.numeric(Weather),
        Surface_Condition_Num = as.numeric(Surface_Condition),
        Light_Num = as.numeric(Light),
        Vehicle_Make_Num = as.numeric(Vehicle_Make),
        Vehicle_Model_Num = as.numeric(Vehicle_Model))

# Perform multiple linear regression
model <- lm(Injury_Severity_Num ~ Weather_Num + Surface_Condition_Num + Light_Num + Vehicle_Year + Vehicle_Make + Vehicle_Model)

# Summary of the regression model
summary(model)

## 
## Call:
## lm(formula = Injury_Severity_Num ~ Weather_Num + Surface_Condition_Num +
##     Light_Num + Vehicle_Year + Vehicle_Make_Num + Vehicle_Model_Num,
##     data = model_data)
## 
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -1.5149 -0.2874 -0.2589 -0.1987  3.0065 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 2.010e+00 1.298e-02 154.778 < 2e-16 ***
## Weather_Num 1.035e-02 9.569e-04 10.814 < 2e-16 ***
## Surface_Condition_Num -7.883e-03 4.148e-04 -19.003 < 2e-16 ***
## Light_Num    -2.456e-03 1.068e-03 -2.300 0.021466 *  
## Vehicle_Year1 -5.815e-02 4.360e-01 -0.133 0.893893  
## Vehicle_Year2 -7.445e-02 6.165e-01 -0.121 0.903883  
## Vehicle_Year3 -3.647e-02 4.360e-01 -0.084 0.933345  
## Vehicle_Year4  5.494e-04 6.165e-01  0.001 0.999289  
## Vehicle_Year8 -4.870e-02 6.165e-01 -0.079 0.937039  
## Vehicle_Year13 -2.182e-02 4.360e-01 -0.050 0.960090  
## Vehicle_Year14 -4.403e-02 6.165e-01 -0.071 0.943058  
## Vehicle_Year15  2.847e-01 3.560e-01  0.800 0.423836  
## Vehicle_Year97  1.003e+00 6.165e-01  1.626 0.103875  
## Vehicle_Year99  6.364e-02 1.213e-01  0.525 0.599920  
## Vehicle_Year198 -2.987e-02 6.165e-01 -0.048 0.961361 

```

## Vehicle_Year200	9.800e-01	4.360e-01	2.248	0.024591	*
## Vehicle_Year201	9.502e-01	4.360e-01	2.179	0.029301	*
## Vehicle_Year202	5.602e-03	6.165e-01	0.009	0.992750	
## Vehicle_Year215	1.953e+00	6.165e-01	3.168	0.001533	**
## Vehicle_Year999	-1.779e-02	6.165e-01	-0.029	0.976973	
## Vehicle_Year1005	-6.479e-02	6.165e-01	-0.105	0.916303	
## Vehicle_Year1008	-3.260e-02	6.165e-01	-0.053	0.957827	
## Vehicle_Year1012	4.015e-02	6.165e-01	0.065	0.948070	
## Vehicle_Year1014	2.015e-02	4.360e-01	0.046	0.963136	
## Vehicle_Year1015	2.004e+00	6.165e-01	3.250	0.001152	**
## Vehicle_Year1025	-1.297e-02	6.165e-01	-0.021	0.983215	
## Vehicle_Year1111	-2.662e-02	3.084e-01	-0.086	0.931196	
## Vehicle_Year1140	-6.953e-02	6.165e-01	-0.113	0.910197	
## Vehicle_Year1234	2.022e-02	4.360e-01	0.046	0.963006	
## Vehicle_Year1900	1.591e-01	7.769e-02	2.048	0.040535	*
## Vehicle_Year1901	-1.741e-02	3.560e-01	-0.049	0.961009	
## Vehicle_Year1911	-7.048e-02	6.165e-01	-0.114	0.908983	
## Vehicle_Year1930	1.863e-03	6.165e-01	0.003	0.997589	
## Vehicle_Year1938	-4.583e-03	6.165e-01	-0.007	0.994069	
## Vehicle_Year1946	9.262e-01	6.165e-01	1.502	0.133008	
## Vehicle_Year1947	3.002e+00	6.165e-01	4.869	1.12e-06	***
## Vehicle_Year1949	1.997e+00	6.165e-01	3.240	0.001196	**
## Vehicle_Year1955	4.901e-01	4.360e-01	1.124	0.260999	
## Vehicle_Year1959	9.509e-01	4.360e-01	2.181	0.029186	*
## Vehicle_Year1960	-3.975e-02	6.165e-01	-0.064	0.948588	
## Vehicle_Year1961	7.014e-02	6.165e-01	0.114	0.909419	
## Vehicle_Year1963	4.713e-01	4.360e-01	1.081	0.279758	
## Vehicle_Year1965	-2.666e-02	4.360e-01	-0.061	0.951250	
## Vehicle_Year1966	-2.231e-02	3.560e-01	-0.063	0.950047	
## Vehicle_Year1967	5.801e-02	6.165e-01	0.094	0.925036	
## Vehicle_Year1968	-6.748e-03	3.560e-01	-0.019	0.984879	
## Vehicle_Year1969	1.748e-01	2.758e-01	0.634	0.526359	
## Vehicle_Year1970	-1.650e-02	3.560e-01	-0.046	0.963029	
## Vehicle_Year1971	2.269e-03	4.360e-01	0.005	0.995847	
## Vehicle_Year1972	4.603e-01	3.084e-01	1.493	0.135550	
## Vehicle_Year1974	-2.637e-02	2.759e-01	-0.096	0.923848	
## Vehicle_Year1975	3.388e-01	3.560e-01	0.952	0.341269	
## Vehicle_Year1976	1.485e+00	4.360e-01	3.407	0.000657	***
## Vehicle_Year1977	9.593e-01	4.360e-01	2.200	0.027779	*
## Vehicle_Year1978	5.829e-01	1.952e-01	2.986	0.002826	**
## Vehicle_Year1979	9.876e-01	3.560e-01	2.774	0.005539	**
## Vehicle_Year1980	2.229e-01	2.057e-01	1.084	0.278477	
## Vehicle_Year1981	4.879e-01	3.084e-01	1.582	0.113606	
## Vehicle_Year1982	7.589e-01	2.057e-01	3.689	0.000226	***
## Vehicle_Year1983	3.847e-01	1.595e-01	2.412	0.015850	*
## Vehicle_Year1984	2.458e-01	1.782e-01	1.379	0.167842	
## Vehicle_Year1985	2.502e-01	1.078e-01	2.321	0.020288	*
## Vehicle_Year1986	3.165e-01	9.924e-02	3.189	0.001429	**
## Vehicle_Year1987	2.623e-01	7.775e-02	3.374	0.000741	***
## Vehicle_Year1988	2.585e-01	6.580e-02	3.930	8.51e-05	***
## Vehicle_Year1989	2.367e-01	5.941e-02	3.985	6.76e-05	***
## Vehicle_Year1990	2.590e-01	5.152e-02	5.027	5.00e-07	***
## Vehicle_Year1991	2.370e-01	5.009e-02	4.732	2.23e-06	***
## Vehicle_Year1992	2.833e-01	4.501e-02	6.295	3.08e-10	***

## Vehicle_Year1993	3.363e-01	3.664e-02	9.178	< 2e-16 ***
## Vehicle_Year1994	3.060e-01	3.116e-02	9.818	< 2e-16 ***
## Vehicle_Year1995	2.639e-01	2.673e-02	9.871	< 2e-16 ***
## Vehicle_Year1996	3.248e-01	2.406e-02	13.498	< 2e-16 ***
## Vehicle_Year1997	2.806e-01	2.007e-02	13.982	< 2e-16 ***
## Vehicle_Year1998	2.841e-01	1.835e-02	15.479	< 2e-16 ***
## Vehicle_Year1999	2.547e-01	1.671e-02	15.247	< 2e-16 ***
## Vehicle_Year2000	2.609e-01	1.485e-02	17.565	< 2e-16 ***
## Vehicle_Year2001	2.573e-01	1.461e-02	17.617	< 2e-16 ***
## Vehicle_Year2002	2.646e-01	1.389e-02	19.043	< 2e-16 ***
## Vehicle_Year2003	2.763e-01	1.323e-02	20.886	< 2e-16 ***
## Vehicle_Year2004	2.655e-01	1.285e-02	20.654	< 2e-16 ***
## Vehicle_Year2005	2.676e-01	1.274e-02	21.009	< 2e-16 ***
## Vehicle_Year2006	2.511e-01	1.245e-02	20.171	< 2e-16 ***
## Vehicle_Year2007	2.483e-01	1.231e-02	20.167	< 2e-16 ***
## Vehicle_Year2008	2.356e-01	1.232e-02	19.124	< 2e-16 ***
## Vehicle_Year2009	2.320e-01	1.268e-02	18.304	< 2e-16 ***
## Vehicle_Year2010	2.507e-01	1.245e-02	20.137	< 2e-16 ***
## Vehicle_Year2011	2.356e-01	1.221e-02	19.291	< 2e-16 ***
## Vehicle_Year2012	2.557e-01	1.203e-02	21.253	< 2e-16 ***
## Vehicle_Year2013	2.448e-01	1.178e-02	20.779	< 2e-16 ***
## Vehicle_Year2014	2.383e-01	1.172e-02	20.338	< 2e-16 ***
## Vehicle_Year2015	2.505e-01	1.163e-02	21.545	< 2e-16 ***
## Vehicle_Year2016	2.324e-01	1.174e-02	19.798	< 2e-16 ***
## Vehicle_Year2017	2.322e-01	1.203e-02	19.295	< 2e-16 ***
## Vehicle_Year2018	2.351e-01	1.263e-02	18.621	< 2e-16 ***
## Vehicle_Year2019	2.268e-01	1.312e-02	17.286	< 2e-16 ***
## Vehicle_Year2020	2.184e-01	1.454e-02	15.019	< 2e-16 ***
## Vehicle_Year2021	2.205e-01	1.572e-02	14.025	< 2e-16 ***
## Vehicle_Year2022	2.149e-01	1.814e-02	11.846	< 2e-16 ***
## Vehicle_Year2023	2.500e-01	2.487e-02	10.050	< 2e-16 ***
## Vehicle_Year2024	2.289e-01	7.770e-02	2.946	0.003222 **
## Vehicle_Year2025	-5.511e-02	6.165e-01	-0.089	0.928774
## Vehicle_Year2033	-2.596e-02	3.560e-01	-0.073	0.941863
## Vehicle_Year2040	-4.317e-03	6.165e-01	-0.007	0.994413
## Vehicle_Year2041	-6.380e-02	6.165e-01	-0.103	0.917580
## Vehicle_Year2048	-1.418e-02	6.165e-01	-0.023	0.981644
## Vehicle_Year2055	-6.124e-02	6.165e-01	-0.099	0.920872
## Vehicle_Year2099	-3.137e-02	6.165e-01	-0.051	0.959421
## Vehicle_Year2100	-4.683e-02	6.165e-01	-0.076	0.939448
## Vehicle_Year2101	-5.775e-02	6.165e-01	-0.094	0.925366
## Vehicle_Year2102	4.464e-02	6.165e-01	0.072	0.942282
## Vehicle_Year2103	-1.125e-03	4.360e-01	-0.003	0.997941
## Vehicle_Year2104	5.026e-01	4.360e-01	1.153	0.249033
## Vehicle_Year2105	1.375e-02	6.165e-01	0.022	0.982211
## Vehicle_Year2107	-4.734e-02	6.165e-01	-0.077	0.938794
## Vehicle_Year2109	2.980e+00	6.165e-01	4.834	1.34e-06 ***
## Vehicle_Year2200	-5.122e-02	6.165e-01	-0.083	0.933788
## Vehicle_Year2201	-2.617e-02	6.165e-01	-0.042	0.966139
## Vehicle_Year2204	-3.336e-02	6.165e-01	-0.054	0.956844
## Vehicle_Year2205	-6.765e-03	6.165e-01	-0.011	0.991246
## Vehicle_Year2208	1.960e+00	6.165e-01	3.180	0.001473 **
## Vehicle_Year2911	4.544e-02	6.165e-01	0.074	0.941247
## Vehicle_Year2912	3.775e-03	4.360e-01	0.009	0.993092

```

## Vehicle_Year2914      -2.984e-02  3.560e-01  -0.084  0.933205
## Vehicle_Year2915      -9.341e-03  6.165e-01  -0.015  0.987911
## Vehicle_Year2916       1.928e-02  4.360e-01   0.044  0.964736
## Vehicle_Year2917      1.004e+00  6.165e-01   1.628  0.103450
## Vehicle_Year2918      -3.586e-02  6.165e-01  -0.058  0.953609
## Vehicle_Year2919      -4.610e-02  6.165e-01  -0.075  0.940393
## Vehicle_Year2991       9.700e-01  6.165e-01   1.573  0.115613
## Vehicle_Year2996       2.825e-01  3.560e-01   0.793  0.427531
## Vehicle_Year2997       6.581e-04  6.165e-01   0.001  0.999148
## Vehicle_Year2998       9.206e-01  6.165e-01   1.493  0.135345
## Vehicle_Year3003       4.915e-01  4.360e-01   1.127  0.259642
## Vehicle_Year3013       3.649e-02  6.165e-01   0.059  0.952806
## Vehicle_Year3863       -6.948e-02 6.165e-01  -0.113  0.910264
## Vehicle_Year5005       -3.093e-02 6.165e-01  -0.050  0.959979
## Vehicle_Year7817       2.504e-02  6.165e-01   0.041  0.967605
## Vehicle_Year8008       -1.784e-02 6.165e-01  -0.029  0.976910
## Vehicle_Year8888       -1.014e-02 4.360e-01  -0.023  0.981442
## Vehicle_Year9999       2.071e-02  8.088e-02   0.256  0.797940
## Vehicle_Make_Num       3.161e-05  2.773e-06  11.397 < 2e-16 ***
## Vehicle_Model_Num      -5.816e-06 7.438e-07  -7.820 5.31e-15 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6164 on 172238 degrees of freedom
## Multiple R-squared:  0.008243,  Adjusted R-squared:  0.007425
## F-statistic: 10.08 on 142 and 172238 DF,  p-value: < 2.2e-16

```

```

# Predictions using the model
predictions <- predict(model)

# Visualize the actual vs. predicted values
ggplot(model_data, aes(x = predictions, y = Injury_Severity_Num)) +
  geom_point() +
  geom_smooth() +
  geom_abline(intercept = 0, slope = 1, color = "red", linetype = "dashed") +
  labs(x = "Predicted Injury Severity", y = "Actual Injury Severity", title = "Actual vs. Predicted Injury Severity")

## 'geom_smooth()' using method = 'gam' and formula = 'y ~ s(x, bs = "cs")'

```

